

산탄데르 고객 만족 예측 - 분류

학습 내용

- 캐글의 산탄데르 고객 만족 데이터 세트에 대해 고객 만족 여부를 XGBoost와 LightGBM을 활용하여 예측

데이터 설명

- 데이터 다운로드 : <https://www.kaggle.com/c/santander-customer-satisfaction/data>
(<https://www.kaggle.com/c/santander-customer-satisfaction/data>)
- 370개의 피처로 이루어진 데이터
- 피처 이름은 전부 익명처리되어 있음.
- 클래스 레이블명은 TARGET
 - 값이 1이면 불만을 가지고 있음.
 - 값이 0이면 만족한 고객

평가 지표

- 성능 평가는 ROC-AUC(ROC 곡선 영역)으로 평가

데이터 로드 및 전처리

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib
```

In [26]:

```
train = pd.read_csv("../data/santander_customer/train.csv", encoding='latin-1')
test = pd.read_csv("../data/santander_customer/test.csv", encoding='latin-1')
sub = pd.read_csv("../data/santander_customer/sample_submission.csv")

train.shape, test.shape, sub.shape
```

Out[26]:

```
((76020, 371), (75818, 370), (75818, 2))
```

In [27]:



```
train.head()
```

Out[27]:

	ID	var3	var15	imp_ent_var16_ult1	imp_op_var39_comer_ult1	imp_op_var39_comer_ult3	im
0	1	2	23	0.0	0.0	0.0	
1	3	2	34	0.0	0.0	0.0	
2	4	2	23	0.0	0.0	0.0	
3	8	2	37	0.0	195.0	195.0	
4	10	2	39	0.0	0.0	0.0	

5 rows × 371 columns

In [28]:



```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 76020 entries, 0 to 76019
Columns: 371 entries, ID to TARGET
dtypes: float64(111), int64(260)
memory usage: 215.2 MB
```

- 111개의 피처가 float형,
- 260개의 피처가 int형
- 모든 피처가 숫자형이며
- NULL값은 없다.

In [29]:



```
train.columns
```

Out[29]:

```
Index(['ID', 'var3', 'var15', 'imp_ent_var16_ult1', 'imp_op_var39_comer_ult1',
      'imp_op_var39_comer_ult3', 'imp_op_var40_comer_ult1',
      'imp_op_var40_comer_ult3', 'imp_op_var40_efect_ult1',
      'imp_op_var40_efect_ult3',
      ...
      'saldo_medio_var33_hace2', 'saldo_medio_var33_hace3',
      'saldo_medio_var33_ult1', 'saldo_medio_var33_ult3',
      'saldo_medio_var44_hace2', 'saldo_medio_var44_hace3',
      'saldo_medio_var44_ult1', 'saldo_medio_var44_ult3', 'var38', 'TARGET'],
      dtype='object', length=371)
```

전체 데이터의 만족(0), 불만족(1) 비율

In [30]:



```
train['TARGET'].value_counts()
```

Out[30]:

```
0    73012
1     3008
Name: TARGET, dtype: int64
```

In [31]:



```
unsatisfied = train['TARGET'].value_counts()[1]
unsatisfied / train['TARGET'].count() # 비율
```

Out[31]:

```
0.0395685345961589
```

In [32]:



```
train.describe()
```

Out[32]:

	ID	var3	var15	imp_ent_var16_ult1	imp_op_var39_comer_u
count	76020.000000	76020.000000	76020.000000	76020.000000	76020.000000
mean	75964.050723	-1523.199277	33.212865	86.208265	72.363000
std	43781.947379	39033.462364	12.956486	1614.757313	339.315800
min	1.000000	-999999.000000	5.000000	0.000000	0.000000
25%	38104.750000	2.000000	23.000000	0.000000	0.000000
50%	76043.000000	2.000000	28.000000	0.000000	0.000000
75%	113748.750000	2.000000	40.000000	0.000000	0.000000
max	151838.000000	238.000000	105.000000	210000.000000	12888.030000

8 rows × 371 columns

- var3의 최소값이 -999999 - 이상치로 보임

In [33]:

```
train['var3'].value_counts()
```

Out[33]:

```
2          74165
8           138
-999999     116
9           110
3           108
```

```
...
218         1
215         1
151         1
87          1
191         1
```

Name: var3, Length: 208, dtype: int64

In [34]:

```
# -999999를 가장 많은 값으로 변경
train['var3'].replace(-999999, 2, inplace=True)
```

In [35]:

```
## ID 열을 삭제
# train.drop('ID', axis=1, inplace=True)
train = train.loc[:, "var3": ]
train.head()
```

Out[35]:

	var3	var15	imp_ent_var16_ult1	imp_op_var39_comer_ult1	imp_op_var39_comer_ult3	imp_o
0	2	23	0.0	0.0	0.0	
1	2	34	0.0	0.0	0.0	
2	2	23	0.0	0.0	0.0	
3	2	37	0.0	195.0	195.0	
4	2	39	0.0	0.0	0.0	

5 rows × 370 columns

In [37]:

```
# 피처와 레이블을 지정.
X = train.iloc[:, :-1]
y = train['TARGET']

X.shape, y.shape
```

Out[37]:

```
((76020, 369), (76020,))
```

In [38]:

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=0)

X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[38]:

```
((60816, 369), (15204, 369), (60816,), (15204,))
```

In [43]:

```
## 레이블 분포비율
print( "학습용 레이블 분포 비율 : Wn" , y_train.value_counts() / y_train.count() )
print( "테스트용 레이블 분포 비율 : Wn" , y_test.value_counts() / y_test.count() )
```

학습용 레이블 분포 비율 :

```
0    0.960964
```

```
1    0.039036
```

Name: TARGET, dtype: float64

테스트용 레이블 분포 비율 :

```
0    0.960964
```

```
1    0.039036
```

Name: TARGET, dtype: float64

모델 생성 및 학습, 그리고 평가해 보기

In [52]:

```
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score

xgb_model = XGBClassifier(n_estimators=500, random_state=156)
xgb_model.fit(X_train, y_train, early_stopping_rounds=100,
              eval_metric='auc', eval_set=[(X_train, y_train), (X_test, y_test)])
```

```
[30] validation_0-auc:0.89741
```

```
[31] validation_0-auc:0.89916
```

```
[32] validation_0-auc:0.90106
```

```
[33] validation_0-auc:0.90253
```

```
[34] validation_0-auc:0.90278
```

```
[35] validation_0-auc:0.90293
```

```
[36] validation_0-auc:0.90463
```

```
[37] validation_0-auc:0.90500
```

```
[38] validation_0-auc:0.90519
```

```
[39] validation_0-auc:0.90533
```

```
[40] validation_0-auc:0.90575
```

```
[41] validation_0-auc:0.90691
```

```
[42] validation_0-auc:0.90716
```

```
[43] validation_0-auc:0.90737
```

```
[44] validation_0-auc:0.90759
```

```
[45] validation_0-auc:0.90769
```

```
[46] validation_0-auc:0.90779
```

```
[47] validation_0-auc:0.90793
```

```
[48] validation_0-auc:0.90831
```

```
[49] validation_0-auc:0.90871
```

```
[50] validation_0-auc:0.90888
```

```
validation_1-auc:0.83335
```

```
validation_1-auc:0.83952
```

```
validation_1-auc:0.83901
```

```
validation_1-auc:0.83885
```

```
validation_1-auc:0.83887
```

```
validation_1-auc:0.83864
```

```
validation_1-auc:0.83834
```

```
validation_1-auc:0.83810
```

```
validation_1-auc:0.83810
```

```
validation_1-auc:0.83813
```

```
validation_1-auc:0.83776
```

```
validation_1-auc:0.83720
```

```
validation_1-auc:0.83684
```

```
validation_1-auc:0.83672
```

```
validation_1-auc:0.83674
```

```
validation_1-auc:0.83693
```

```
validation_1-auc:0.83686
```

```
validation_1-auc:0.83678
```

```
validation_1-auc:0.83694
```

```
validation_1-auc:0.83676
```

```
validation_1-auc:0.83655
```

In [53]:



```
pred_prob = xgb_model.predict_proba(X_test)[:, 1]
pred_prob
```

Out [53]:

```
array([0.00643863, 0.02387667, 0.01260844, ..., 0.05883254, 0.01729385,
       0.01727541], dtype=float32)
```

In [54]:



```
xgb_roc_score = roc_auc_score(y_test, pred_prob, average='macro')
print("ROC AUC : {:.4f}".format(xgb_roc_score))
```

ROC AUC : 0.8413

하이퍼 파라미터 튜닝

- max_depth, min_child_weight, colsample_bytree
- 먼저 2-3개 정도의 파라미터를 최적화 시킨 후, 최적 파라미터를 기반으로 1-2개 파라미터를 결합하여 튜닝을 수행

In [55]:



```
%%time

from sklearn.model_selection import GridSearchCV

# 우선 하이퍼 파라미터 수행 속도를 향상을 위해 100으로
xgb_model1 = XGBClassifier(n_estimators=100, use_label_encoder=False)
params = {"max_depth": [5, 7],
          "min_child_weight": [1, 3],
          "colsample_bytree": [0.5, 0.75]}

gridcv = GridSearchCV(xgb_model1, param_grid=params, cv=3)
gridcv.fit(X_train, y_train, early_stopping_rounds=30,
          eval_metric='auc',
          eval_set = [(X_train, y_train), (X_test, y_test)])

[40] validation_0-auc:0.87678      validation_1-auc:0.83859
[41] validation_0-auc:0.87711      validation_1-auc:0.83830
[42] validation_0-auc:0.87738      validation_1-auc:0.83823
[43] validation_0-auc:0.87752      validation_1-auc:0.83796
[44] validation_0-auc:0.87777      validation_1-auc:0.83765
[45] validation_0-auc:0.87785      validation_1-auc:0.83786
[46] validation_0-auc:0.87802      validation_1-auc:0.83761
[47] validation_0-auc:0.87840      validation_1-auc:0.83698
[48] validation_0-auc:0.87868      validation_1-auc:0.83699
[49] validation_0-auc:0.87882      validation_1-auc:0.83708
[0] validation_0-auc:0.80039      validation_1-auc:0.80013
[1] validation_0-auc:0.82111      validation_1-auc:0.82026
[2] validation_0-auc:0.82749      validation_1-auc:0.82627
[3] validation_0-auc:0.83124      validation_1-auc:0.82830
[4] validation_0-auc:0.83475      validation_1-auc:0.82881
[5] validation_0-auc:0.83676      validation_1-auc:0.83385
[6] validation_0-auc:0.83648      validation_1-auc:0.83085
[7] validation_0-auc:0.84336      validation_1-auc:0.83472
[8] validation_0-auc:0.84624      validation_1-auc:0.83404
[9] validation_0-auc:0.84541      validation_1-auc:0.83287
```

In [64]:



```
print("GridSearchCV 최적 파라미터 :", gridcv.best_params_)

pred_prob = gridcv.predict_proba(X_test)[: , 1]
xgb_roc_score = roc_auc_score(y_test, pred_prob, average='macro')
print("ROC AUC : {0:4f}".format(xgb_roc_score))
```

GridSearchCV 최적 파라미터 : {'colsample_bytree': 0.5, 'max_depth': 5, 'min_child_weight': 3}
ROC AUC : 0.844455

실습해 보기

- colsample_bytree : 0.5, max_depth : 5, min_child_weight : 3로 설정
- n_estimators = 1000으로 증가, learning_rate를 조정해보고, reg_alpha를 추가하여 ROC_AUC의 값을 구해보자.

In [59]:



```
%%time

xgb_model_l = XGBClassifier(n_estimators=1000,
                           random_state= 77,
                           learning_rate=0.02,
                           max_depth=5,
                           min_child_weight=3,
                           colsample_bytree=0.5,
                           reg_alpha=0.03)

# 성능 평가 지표를 auc로, 조기 중단 파라미터 값은 200으로 설정하고 학습 수행
xgb_model_l.fit(X_train, y_train, early_stopping_rounds=200,
               eval_metric='auc', eval_set=[(X_train, y_train), (X_test, y_test)])
```

```
[525] validation_0-auc:0.88177      validation_1-auc:0.84484
[526] validation_0-auc:0.88183      validation_1-auc:0.84480
[527] validation_0-auc:0.88188      validation_1-auc:0.84481
[528] validation_0-auc:0.88192      validation_1-auc:0.84479
[529] validation_0-auc:0.88197      validation_1-auc:0.84479
[530] validation_0-auc:0.88200      validation_1-auc:0.84481
[531] validation_0-auc:0.88205      validation_1-auc:0.84484
[532] validation_0-auc:0.88213      validation_1-auc:0.84478
Wall time: 1min 48s
```

Out [59]:

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=0.5, gamma=0, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.02, max_delta_step=0, max_depth=5,
              min_child_weight=3, missing=nan, monotone_constraints='()',
              n_estimators=1000, n_jobs=8, num_parallel_tree=1, random_state=77,
              reg_alpha=0.03, reg_lambda=1, scale_pos_weight=1, subsample=1,
              tree_method='exact', validate_parameters=1, verbosity=None)
```

In [60]:



```
pred_prob = xgb_model_l.predict_proba(X_test)[: , 1]
xgb_roc_score = roc_auc_score(y_test, pred_prob, average='macro')
print("ROC AUC : {0:4f}".format(xgb_roc_score))
```

ROC AUC : 0.845269

메모

- XGBoost는 GBM을 기반으로 하고 있기에, 수행시간이 어느정도 걸립니다.
- 앙상블 계열 알고리즘에서 하이퍼 파라미터 튜닝으로 성능 수치 개선이 급격하게 되는 경우는 많지 않습니다.

각 특징의 중요도 시각화

- xgboost 모듈의 plot_importance() 메서드를 이용

In [63]:



```
from xgboost import plot_importance
import matplotlib.pyplot as plt

fig, ax = plt.subplots(1,1, figsize=(10,8))
plot_importance(xgb_model_1, ax=ax, max_num_features=20, height=0.4)
```

Out[63]:

<AxesSubplot:title={'center':'Feature importance'}, xlabel='F score', ylabel='Features'>

